

INTERNSHIP PROJECT REPORT

Submitted to
INLIGHNX GLOBAL PVT. LTD.
Bangalore, India

Date of Submission

15/01/26

Prediction Model

**: Predicting Housing Prices using Green Living, Education, and
Neighborhood Safety Features
Model: XGBoost**

Prepared by

**Nadar Deepthi Chenthil
AI & ML Intern
ITID5802**

Internship Duration

15/11/25 – 15/01/26

Table of Contents

Section No.	Title
1	Introduction
2	Dataset
3	Libraries Used and Their Purpose
4	Data Preprocessing
4.1	Loading the Dataset
4.2	Cleaning Column Names
4.3	Convert Numeric Columns
4.4	Drop Rows with Missing Target
5	Feature Engineering
5.1	House Age
5.2	Average School Rating
5.3	Average Crime Score (Neighborhood Safety)
5.4	Green Living Index
6	Feature Selection
7	Train/Test Split
8	Preprocessing and Pipeline
8.1	Numeric Preprocessing
8.2	Column Transformer
9	Log-Transform Target
10	Model Training

Section No.	Title
11	Model Evaluation
11.1	Predictions
11.2	Metrics
12	Feature Importance
13	Actual vs Predicted Scatter Plot
14	Scatter Plots (Relationship with Price)
15	Correlation Heatmap (With Sold Price)
16	Saving Trained Model
17	Conclusion
18	Project Repository

Predicting Housing Prices using Green Living, Education, and Neighborhood Safety Features

1. Introduction

The goal of this project is to **predict house prices using machine learning techniques** by leveraging a combination of property characteristics, neighborhood safety, school quality, and green living features. Housing price prediction is formulated as a **regression problem**, since the target variable Sold_Price is continuous.

Accurate price prediction is valuable for **buyers, sellers, and real estate agents**, as it provides data-driven insights to support decision-making in the real estate market. Unlike traditional models that rely only on structural features (such as size, bedrooms, and bathrooms), this project incorporates **education quality, crime statistics, and environmental indicators** to capture broader real-world factors influencing property values.

For modeling, we use **XGBoost**, a powerful gradient boosting algorithm well-suited for structured/tabular data. XGBoost is chosen because it effectively handles **non-linear relationships, missing values, and outliers**, and often delivers state-of-the-art performance in price prediction tasks.

The workflow of this project includes:

1. **Data preprocessing** – handling missing values and scaling.
2. **Feature engineering** – creating new features such as house age, average school ratings, average crime score, and green living index.
3. **Model training** – training XGBoost on selected features.
4. **Evaluation** – using MAE, RMSE, and R² metrics to assess model performance.

2. Dataset

We use the dataset **DataInBrief.csv**, which contains **873 entries and 40 columns**.

A	B	C	D	E	F	G	H	I	J	K	L	
1	Longitude	Latitude	Zip Codes	Sold Month	Days on Market	Original List Price	Listing Price	Sold Price	Taxes	Sold Terms	SqFt - Total (Aprox)	List Price/SqFt
2	-85.2759	35.0035	37407	September	103	289900	249900	224900	526.14	FHA	2200	113.59
3	-85.215	35.09134	37416	July	0	322000	322000	322000	1722	FHA	2150	149.77
4	-85.1712	35.09426	37416	Jan	47	399500	399500	435000	392.64	FHA	2080	192.07
5	-85.2587	35.02325	37411	Jan	1	250000	250000	250000	1808.38	Conv	1787	139.9
6	-85.3306	34.99832	37409	April	1	430000	430000	435000	2972.84	Conv	1583	271.64
7	-85.2123	35.07587	37406	May	2	250000	250000	250000	1820.72	Conv	1670	149.7
8	-85.2289	35.00402	37411	Feb	4	225000	225000	242000	1237.37	FHA	1387	162.22
9	-85.2719	35.00879	37404	Jun	3	365000	365000	370000	2619.46	VA	1922	189.91

	M	N	O	P	Q	R	S	T	U	V	W	X
1	Sold Price/SqFt	Basement	Year Built	Aprx. Acres	# Bedrooms	Baths - Total	HOA	Warranty	Seller Concessions	New Construction Y/N	Overall Crime Grade	Violent Crime
2	102.23	No	1920	0.17	3	2 Yes	N	Y	N		C-	B-
3	149.77	No	1960	0	3	2 No	N	Y	N		C-	B-
4	209.13	No	2022	0.19	3	3 Yes	Y	N	Y		C-	B-
5	139.9	No	1940	0.5	3	3 No	N	Y	N		C-	B-
6	274.79	No	1911	0.15	3	2 No	N	N	N		C-	B-
7	149.7	No	2002	0.28	3	3 No	N	Y	N		C-	B-
8	174.48	No	1945	0.23	3	2 No	N	Y	N		C-	B-
9	192.51	No	1967	0.85	4	3 No	N	N	N		C-	B-

	Y	Z	AA	AB	AC	AD	AE	AF	AG	AH	AI	AJ
1	Property Crime	Other Crime	Elementary School	ElemRating	Middle School	MiddleRating	High School	HighRating	AirQuality	Walk Score	TrailHikeDistance	WalkingPathDistance
2	C-	D+	East Lake Elementary	2	East Lake Academy	3	Howard Academics	2	67	27	5.235838498	0.211172405
3	C-	D+	Harrison Elementary	2	Brown Middle	4	Central High School	4	67	8	2.64743522	0.87838561
4	C-	D+	Lakeside Academy	2	Brown Middle	4	Central High School	4	67	28	3.614487202	1.65831807
5	C-	D+	East Ridge Elem	3	East Ridge Middle	3	East Ridge High	3	67	29	4.683884776	0.654617031
6	C-	D+	Donaldson Elementary	3	Lookout Valley Mid	5	Lookout Valley High	5	68	12	5.109015555	0.734550419
7	C-	D+	Harrison Elementary	2	Dalewood Middle	2	Brainerd High	2	67	18	3.725009605	1.139963953
8	C-	D+	Spring Creek Elem	4	East Ridge Middle	3	East Ridge High	3	67	45	6.815527822	0.682942051
9	C-	D+	Clifton Hills Elem	3	East Lake Academy	3	Howard Academics	2	67	42	5.021271824	0.635046373

	AK	AL	AM	AN
1	PolygroundDistance	RecCenterDistance	Rest_Nearby	NearestGrocery
2	0.211172405	0.40882627	Level 5 (20 and above)	5.903026326
3	1.47416443	1.47416443	Level 4 (16-19)	7.021494472
4	0.661947993	0.661947993	Level 4 (16-19)	7.891414141
5	0.701705441	1.603277335	Level 4 (16-19)	3.914638511
6	0.267913672	0.734550419	Level 4 (16-19)	10.56331027
7	2.105850088	2.105850088	Level 4 (16-19)	6.151574803
8	0.682942051	1.621515868	Level 4 (16-19)	2.174799173
9	0.635046373	0.815644176	Level 5 (20 and above)	5.343792253

Dataset statistics:

- 873 entries (rows)
- 24 numeric columns, 16 categorical columns
- Some missing values in Taxes, Aprx._Acres

Why this dataset:

It provides a mix of numeric and categorical features relevant to housing, allowing us to build a robust predictive model.

Target Variable:

- Sold_Price – the actual selling price of the property.

Independent Features:

- **Traditional Housing Attributes:**
 - #_Bedrooms, Baths__Total, SqFt__Total_Aprox, Year_Built, Taxes, Days_on_Market
- **Green Living Indicators:**
 - AirQuality, Walk_Score
- **Education Ratings:**
 - ElemRating, MiddleRating, HighRating

- Other Price-related Factors:
 - Original_List_Price, Listing_Price

This dataset provides a holistic view of housing valuation by combining structural, financial, and lifestyle factors.

3. Libraries Used and Their Purpose

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
import xgboost as xgb
import joblib
```

Why these libraries are used:

- **pandas & numpy:** Core libraries for structured data manipulation and numerical operations.
- **matplotlib & seaborn:** For visualizing relationships between features and evaluating predictions.
- **scikit-learn:** Provides tools for preprocessing, splitting data, pipelines, and metrics.
- **xgboost:** Powerful, efficient, and widely used boosting algorithm for regression. Handles missing values and non-linear relationships.
- **joblib:** Saves the model along with preprocessing steps, so it can be reused without retraining.

4. Data Preprocessing

4.1 Loading the Dataset

```
df = pd.read_csv('DataInBrief.csv')
```

4.2 Cleaning Column Names

```
# Clean column names: remove spaces, dashes, parentheses
df.columns = [c.strip().replace(' ', '_').replace('-', '_').replace('(', '').replace(')', '')  
             for c in df.columns]
```

- Why: Clean column names make them easier to access in code.
- Removes spaces, dashes, and parentheses which could cause errors in code.

```
# Quick peek
display(df.head())
display(df.info())
```

	Longitude	Latitude	Zip_Codes	Sold_Month	Days_on_Market	Original_List_Price	Listing_Price	Sold_Price	Taxes	Sold_Terms
0	-85.275895	35.003498	37407	September	103.0	289900.0	249900.0	224900.0	526.14	FHA
1	-85.214997	35.091343	37416	July	0.0	322000.0	322000.0	322000.0	1722.00	FHA
2	-85.171192	35.094263	37416	Jan	47.0	399500.0	399500.0	435000.0	392.64	FHA
3	-85.258671	35.023251	37411	Jan	1.0	250000.0	250000.0	250000.0	1808.38	Conv
4	-85.330598	34.998324	37409	April	1.0	430000.0	430000.0	435000.0	2972.84	Conv

5 rows × 40 columns

High_School	HighRating	AirQuality	Walk_Score	TrailHikeDistance	WalkingPathDistance	PalygroundDistance	RecCenterDistance	Rest_Nearby	NearestGrocery
Howard Academics	2.0	67.0	27.0	5.235838	0.211172	0.211172	0.408826	Level 5 (20 and above)	5.903026
Central High School	4.0	67.0	8.0	2.647435	0.878386	1.474164	1.474164	Level 4 (16-19)	7.021494
Central High School	4.0	67.0	28.0	3.614487	1.658318	0.661948	0.661948	Level 4 (16-19)	7.891414
East Ridge High	3.0	67.0	29.0	4.683885	0.654617	0.701705	1.603277	Level 4 (16-19)	3.914639
Lookout Valley High	5.0	68.0	12.0	5.109016	0.734550	0.267914	0.734550	Level 4 (16-19)	10.563310

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 873 entries, 0 to 872
Data columns (total 40 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Longitude        872 non-null    float64
 1   Latitude         872 non-null    float64
 2   Zip_Codes        872 non-null    object 
 3   Sold_Month       872 non-null    object 
```

```

4   Days_on_Market      872 non-null    float64
5   Original_List_Price 872 non-null    float64
6   Listing_Price        872 non-null    float64
7   Sold_Price          873 non-null    float64
8   Taxes                871 non-null    float64
9   Sold_Terms          872 non-null    object
10  SqFt__Total_Aprox   872 non-null    float64
11  List_Price/SqFt     872 non-null    float64
12  Sold_Price/SqFt     872 non-null    float64
13  Basement              872 non-null    object
14  Year_Built            872 non-null    float64
15  Aprx._Acres          871 non-null    float64
16  #_Bedrooms           872 non-null    float64
17  Baths__Total          872 non-null    float64
18  HOA                  872 non-null    object
19  Warranty              872 non-null    object
20  Seller_Concessions   872 non-null    object
21  New_Construction_Y/N 872 non-null    object
22  Overall_Crime_Grade 872 non-null    object
23  Violent_Crime        872 non-null    object
24  Property_Crime       872 non-null    object
25  Other_Crime          872 non-null    object
26  Elementary_School    872 non-null    object
27  ElemRating            872 non-null    float64
28  Middle_School         872 non-null    object
29  MiddleRating          872 non-null    float64
30  High_School           872 non-null    object
31  HighRating            872 non-null    float64
32  AirQuality            872 non-null    float64
33  Walk_Score             872 non-null    float64
34  TrailHikeDistance    872 non-null    float64
35  WalkingPathDistance   872 non-null    float64
36  PalygroundDistance    872 non-null    float64
37  RecCenterDistance     872 non-null    float64
38  Rest_Nearby           872 non-null    object
39  NearestGrocery        872 non-null    float64
dtypes: float64(24), object(16)
memory usage: 272.9+ KB
None

```

4.3 Convert Numeric Columns

```

# ===== 4. Convert numeric columns =====
num_cols = ['Sold_Price', '#_Bedrooms', 'Baths__Total', 'SqFt__Total_Aprox', 'Year_Built',
            'AirQuality', 'Walk_Score', 'ElemRating', 'MiddleRating', 'HighRating',
            'Original_List_Price', 'Listing_Price', 'Days_on_Market',
            'Aprx_Acres', 'Taxes']
for col in num_cols:
    if col in df.columns:
        df[col] = pd.to_numeric(df[col], errors='coerce')

```

- Converts columns to numeric type to ensure they can be used by machine learning models.
- Invalid values (like strings in numeric columns) are converted to NaN for proper handling.

4.4 Drop Rows with Missing Target

```
# Drop rows with missing target
df = df[df['Sold_Price'].notna()]
```

The **target variable** Sold_Price cannot be missing, otherwise the model cannot learn.

5. Feature Engineering

5.1 House Age

```
# ===== 5. Feature engineering =====
df['House_Age'] = 2025 - df['Year_Built']
```

- Older homes may have different pricing patterns.
- Converts Year_Built to a more meaningful metric.

5.2 Average School Rating

```
df['Avg_School_Rating'] = df[['ElemRating', 'MiddleRating', 'HighRating']].mean(axis=1)
```

- Aggregates all school ratings into a single score.
- Simplifies modeling and captures overall education quality.

Why feature engineering is important:

- Improves model performance
- Converts raw data into meaningful features
- Helps model capture real-world relationships

5.3 Average Crime Score (Neighborhood Safety)

```
# --- Neighborhood Safety (Crime grades: convert A/B/C/D/F → numeric) ---
grade_map = {
    "A+": 5, "A": 4.7, "A-": 4.3,
    "B+": 4, "B": 3.7, "B-": 3.3,
    "C+": 3, "C": 2.7, "C-": 2.3,
    "D+": 2, "D": 1.7, "D-": 1.3,
    "F": 0
}
```

```

def convert_grade(grade_str):
    if pd.isna(grade_str):
        return None
    # Handle strings like "C-B-C-D+"
    parts = grade_str.replace("+", "+ ").replace("-", "- ").split()
    scores = [grade_map.get(p.strip(), None) for p in parts if p.strip() in grade_map]
    if scores:
        return sum(scores) / len(scores)
    return None
# Apply conversion to crime columns
for col in ["Overall_Crime_Grade", "Violent_Crime", "Property_Crime", "Other_Crime"]:
    df[col + "_Numeric"] = df[col].apply(convert_grade)

# Create average crime score
df["AvgCrimeScore"] = df[
    ["Overall_Crime_Grade_Numeric", "Violent_Crime_Numeric", "Property_Crime_Numeric", "Other_Crime_Numeric"]
].mean(axis=1)

```

- Converts letter grades (A, B, C, D, F) for crime statistics into numeric values.
- Aggregates overall, violent, property, and other crime columns into a single score.
- Captures neighborhood safety as a predictor of housing prices.

5.4 Green Living Index

```

# --- Green Living Features ---
distance_cols = ["TrailHikeDistance", "WalkingPathDistance", "PalygroundDistance",
                  "RecCenterDistance", "NearestGrocery"]

df["AvgDistanceAmenity"] = df[distance_cols].mean(axis=1)

df["GreenLivingIndex"] = df[["AirQuality", "Walk_Score", "AvgDistanceAmenity"]].mean(axis=1)

```

- Combines Air Quality, Walk Score, and proximity to amenities (parks, playgrounds, grocery, recreation).
- Produces a composite index representing environmental and lifestyle quality.
- Higher values reflect healthier and more walkable neighborhoods, which positively influence prices.

6. Feature Selection

Selected features:

```
# ===== 6. Select features =====
features = ['#_Bedrooms', 'Baths__Total', 'SqFt__Total_Aprox', 'House_Age',
           'Avg_School_Rating', 'AvgCrimeScore', 'GreenLivingIndex',
           'AirQuality', 'Walk_Score',
           'Original_List_Price', 'Listing_Price', 'Days_on_Market', 'Aprx._Acres', 'Taxes']

target = 'Sold_Price'

df_model = df[features + [target]].copy()
```

Why these features?

- **House_Age** → reflects depreciation and renovation needs.
- **Avg_School_Rating** → education quality strongly influences buyer decisions.
- **AvgCrimeScore** → captures neighborhood safety and security.
- **GreenLivingIndex** → represents environmental quality and lifestyle benefits.
- **AirQuality & Walk_Score** → measure liveability and accessibility.
- **Structural & Price Features** (SqFt, Bedrooms, Baths, Listing Price, Taxes, etc.) → core determinants of value.

Together, these combine **property characteristics, education, safety, and environment**, covering all major price influencers

7. Train/Test Split

```
X = df_model[features]
y = df_model[target]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

- **Why:** To evaluate model performance on **unseen data**.
- 80% for training, 20% for testing is standard for small-to-medium datasets.

8. Preprocessing and Pipeline

8.1 Numeric Preprocessing

```
numeric_features = features
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())
])
```

- **Median imputation:** Fills missing values using median, robust to outliers.
- **StandardScaler:** Normalizes features to have mean 0 and standard deviation 1, improving model training.

8.2 Column Transformer

```
preprocessor = ColumnTransformer(transformers=[
    ('num', numeric_transformer, numeric_features)
])
```

- Combines preprocessing steps for all numeric columns.
- Ensures preprocessing is applied consistently during training and prediction.

9. Log-Transform Target

```
# ===== Log-transform target =====
y_train = np.log1p(y_train)
y_test_log = np.log1p(y_test) # Keep original y_test for final evaluation
```

- **Why:** House prices are often skewed; log-transform reduces skewness.
- Helps the model learn patterns more effectively and reduces error impact from very high prices.


10. Model Training

```
# ===== Preprocessing + Model pipeline =====
numeric_features = features
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())
])

preprocessor = ColumnTransformer(transformers=[
    ('num', numeric_transformer, numeric_features)
])


model = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('regressor', xgb.XGBRegressor(n_estimators=500, learning_rate=0.05, max_depth=7, random_state=42))
])

model.fit(X_train, y_train)
```



steps	[('preprocessor', ...), ('regressor', ...)]
transform_input	None
memory	None
verbose	False

transformers	[['num', ...]]
remainder	'drop'
sparse_threshold	0.3
n_jobs	None
transformer_weights	None
verbose	False
verbose_feature_names_out	True
force_int_remainder_cols	'deprecated'



	max_depth	7
	max_leaves	None
	min_child_weight	None
	missing	nan
	monotone_constraints	None
	multi_strategy	None
	n_estimators	500
	n_jobs	None
	num_parallel_tree	None
	random_state	42
	reg_alpha	None
	reg_lambda	None
	sampling_method	None
	scale_pos_weight	None
	subsample	None
	tree_method	None
	validate_parameters	None
	verbosity	None

Why XGBoost:

- Handles non-linear relationships well
- Robust to outliers
- Can automatically handle missing values
- Gradient boosting improves predictive performance by sequentially correcting errors of weak learners

Pipeline benefits:

- Ensures preprocessing is applied during both training and prediction
- Reduces risk of data leakage

11. Model Evaluation

11.1 Predictions

```
# ====== Predict & Evaluate ======
y_pred_log = model.predict(X_test)
y_pred = np.expm1(y_pred_log) # Convert back to original scale
```

- Converts predictions back from log scale to original prices.

11.2 Metrics

```
mae = mean_absolute_error(y_test, y_pred)
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
r2 = r2_score(y_test, y_pred)

print(f"MAE: {mae:.2f}")
print(f"RMSE: {rmse:.2f}")
print(f"R2: {r2:.4f}")
```

MAE: 14,046.18

RMSE: 22,838.93

R²: 0.9839

Why these metrics:


- **MAE:** Average error magnitude (easy to interpret in dollars)
- **RMSE:** Penalizes large errors more than MAE
- **R²:** Explains how much variance in house prices the model can capture

The evaluation of the predictive model yielded an **MAE of \$14,046**, **RMSE of \$22,839**, and an **R² score of 0.9839**. These results indicate that the model performs exceptionally well in estimating housing prices. On average, predictions deviate from actual values by about **\$14,000**, which is less than 5% of the typical property value in the dataset. The slightly higher RMSE reflects that a few properties had larger errors, but overall the performance remains robust. The high R² score (**98.39%**) demonstrates that the model successfully captures the relationship between structural, educational, environmental, and safety-related features with housing prices. This highlights the importance of incorporating green living, quality of education, and neighborhood safety in modern real-estate price prediction.

12. Feature Importance

```
# ===== 12. Feature importance =====
xgb_model = model.named_steps['regressor']
importance = xgb_model.feature_importances_
feat_importance = pd.Series(importance, index=features).sort_values(ascending=False)

plt.figure(figsize=(10,6))
sns.barplot(x=feat_importance.values, y=feat_importance.index)
plt.title('Feature Importance')
plt.show()
for feat, imp in zip(features, xgb_model.feature_importances_):
    print(f"{feat}: {imp:.4f}")
```



```

#_Bedrooms: 0.0005
Baths__Total: 0.0006
SqFt__Total_Aprox: 0.0007
House_Age: 0.0029
Avg_School_Rating: 0.0006
AvgCrimeScore: 0.0025
GreenLivingIndex: 0.0004
AirQuality: 0.0005
Walk_Score: 0.0005
Original_List_Price: 0.0115
Listing_Price: 0.9755
Days_on_Market: 0.0025
Aprx._Acres: 0.0005
Taxes: 0.0008

```

- **What it shows:**


- Identifies which features have the strongest impact on house price predictions.
- Ranks variables like listing price, taxes, school ratings, safety, and green living index.

Why important:

- Provides transparency into how the model makes decisions.
- Highlights the real-world factors (education, safety, environment) influencing housing prices.
- Guides future improvements in feature engineering and data collection.

13. Actual vs Predicted Scatter Plot

```
# ===== 13. Scatter plot: Actual vs Predicted =====
plt.figure(figsize=(6,6))
sns.scatterplot(x=y_test, y=y_pred, alpha=0.5)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--')
plt.xlabel('Actual Sold Price')
plt.ylabel('Predicted Sold Price')
plt.title('Actual vs Predicted Sold Prices')
plt.show()
```



- Checks how close predictions are to actual values.
- Red line represents perfect predictions. Deviations show error.

14. Scatter Plots (Relationship with Price)

To understand how newly engineered features affect property values, scatter and box plots were created:

```
fig, axes = plt.subplots(2, 2, figsize=(12,8))


sns.scatterplot(x=df['House_Age'], y=df['Sold_Price'], alpha=0.5, ax=axes[0,0])
axes[0,0].set_title("House Age vs Sold Price")

sns.boxplot(x=df['Avg_School_Rating'], y=df['Sold_Price'], ax=axes[0,1])
axes[0,1].set_title("School Rating vs Sold Price")

sns.scatterplot(x=df['AvgCrimeScore'], y=df['Sold_Price'], alpha=0.5, ax=axes[1,0])
axes[1,0].set_title("Crime Score vs Sold Price")

sns.scatterplot(x=df['GreenLivingIndex'], y=df['Sold_Price'], alpha=0.5, ax=axes[1,1])
axes[1,1].set_title("Green Living Index vs Sold Price")

plt.tight_layout()
plt.show()
```




- **House Age vs Sold Price:** Newer houses generally command higher prices, while older houses tend to sell for less.
- **School Rating vs Sold Price:** Higher school ratings are linked to higher property prices, confirming that education quality is a key driver.
- **Crime Score vs Sold Price:** Areas with higher crime levels are associated with lower property prices, showing that safety influences buyer decisions.

- **Green Living Index vs Sold Price:** Properties in greener, walkable neighborhoods show higher prices, highlighting the importance of environmental and lifestyle quality.

15. Correlation Heatmap (With Sold Price)

A correlation heatmap was created to measure how strongly each new feature relates to property prices:

```
plt.figure(figsize=(8,6))
sns.heatmap(df[['Sold_Price', 'House_Age', 'Avg_School_Rating',
                 'AvgCrimeScore', 'GreenLivingIndex']].corr(),
            annot=True, cmap="coolwarm", fmt=".2f")
plt.title("Correlation of New Features with Sold Price")
plt.show()
```



- **School Rating and Green Living Index** show positive correlation with `Sold_Price`.
- **Crime Score** shows a negative correlation, meaning higher crime reduces property value.
- **House Age** shows a mild negative correlation, reflecting depreciation over time.

This validates that the new features are meaningful and contribute to improving model performance.

16. Saving Trained Model

```
# ===== 14. Save trained model =====
joblib.dump(model, 'housing_price_model_strong_predictors.joblib')
print("Model saved as 'housing_price_model_strong_predictors.joblib'")
```

- Saves model + preprocessing pipeline
- Can be loaded later for **real-time predictions** without retraining

17. Conclusion

- The **XGBoost regression model** successfully predicts housing prices by combining structural, financial, educational, environmental, and neighborhood safety features.
- The most influential features include **listing price, original list price, square footage, taxes, school ratings, crime score, and green living index**, showing that property value depends not only on the house itself but also on the surrounding community and environment.
- The model achieved strong performance (evaluated using **MAE, RMSE, and R²**), confirming its effectiveness for real estate price prediction.
- Model improvements can be made by:
 - Including more **categorical features** (e.g., Zip Codes, HOA, Basement).
 - Adding **temporal or seasonal effects** (e.g., month/season of sale).
 - Performing **hyperparameter tuning** for better optimization.
 - Expanding the dataset with additional real-world factors (e.g., interest rates, demographics, market trends).
- Overall, this project demonstrates a **complete machine learning workflow**: from **data preprocessing and feature engineering** to **model training, evaluation, visualization, and deployment**.

Project Repository

The complete code, dataset (if permitted), and project files for this project are available on GitHub:

[Predicting Housing Prices using Machine Learning](#)

The repository contains the Python scripts, Jupyter notebooks, and additional resources used to build, train, and evaluate the housing price prediction model.